

The background is a dark gray field filled with a complex geometric pattern of triangles and diamonds. The shapes are outlined in black and filled with various colors: red, orange, green, blue, and purple. The pattern is symmetrical and creates a sense of depth and movement.

Noms

**A new database for the Big
Data World**

The Developers Conference

Porto Alegre 2016

Big Data

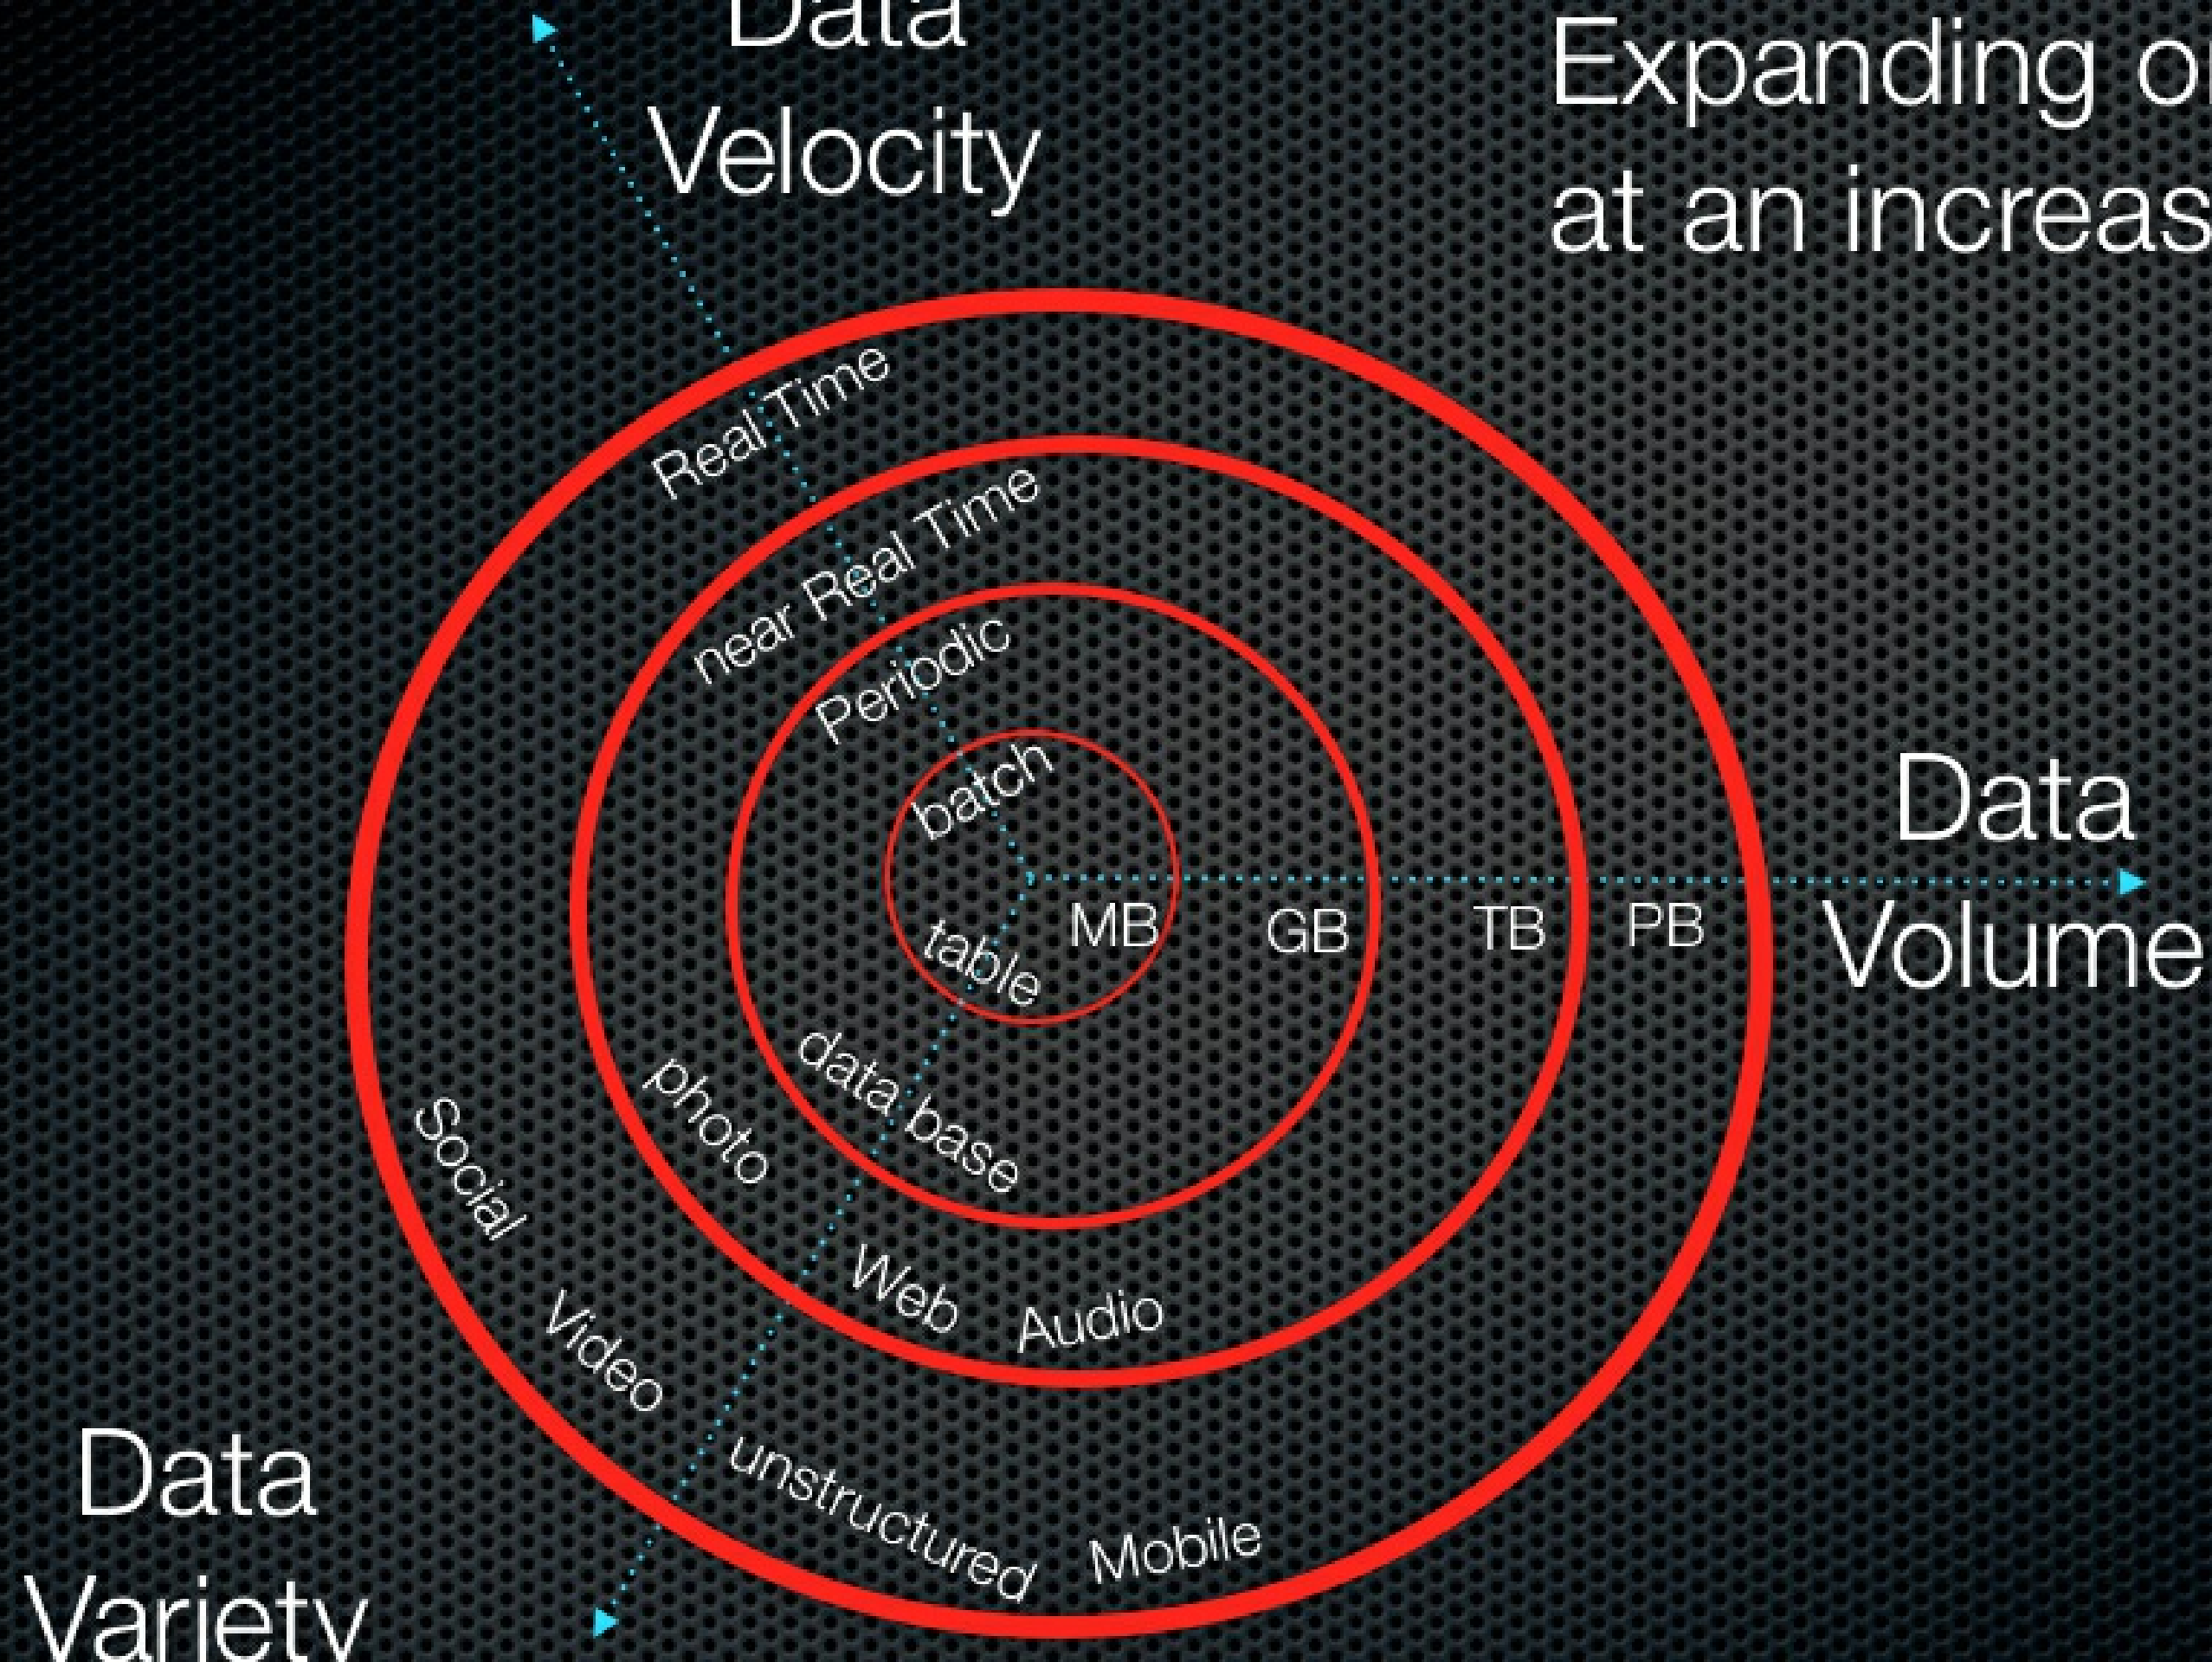
“Big data is a term for data sets that are so large or complex that traditional data processing applications are inadequate to deal with them.”

source: https://en.wikipedia.org/wiki/Big_data

3D Data Management: Controlling Data Volume, Velocity, and Variety

Doug Laney, 6th February 2001

Expanding on 3 fronts
at an increasing rate.



today's challenges include...

sharing & storage

Introducing Noms

"Noms is a decentralized database based on ideas from Git."



Attic Labs

- Founders (Ex-Googlers): **Rafael Weinstein** / **Aaron Boodman**
- Raised \$8.1M (mainly from Greylock Partners)
- Source code available **@GitHub**, developed in Go



About Noms

Decentralized

Data is modeled as a directed acyclic graph

Users can have their own copy of the database

Any peer is free to move forward independently

About Noms

Content-addressed

Intended to store data that does not change over time

Concept of content identifier

Duplicated data is automatically deduplicated

About Noms

Append-only

Data is never removed by default

About Noms

Strongly Typed

Every value has a type

Types are evaluated when you insert new data

Noms X Traditional Database

Noms	Traditional
Any peer's state is valid	Single source of truth
Historical records	Single point in time

Noms Concepts

- Databases / Datasets
- Chunks
- Types
 - Boolean, Number, List, Blob, etc
 - Allow data to be self-describing
 - Every chunk has a header that describes its type
 - Types can be verified by clients

Type Accretion

Widen (or narrow) the type of any container

```
Set<struct {name: String}>
```

becomes

```
Set<struct {name: String}> | struct {name: String, age: Number}>>
```

Project Status

Data Format

Core data format defined

Project Status

Completeness

We're just getting started

Project Status

API

Public API will continue to evolve

Noms CLI & SDK

An introduction to Noms CLI and JS SDK

Installing

Command tools can be downloaded from GitHub or build from source code (requires Go 1.6+)

<https://github.com/attic-labs/noms>

Serving a database

Databases uses Google LevelDB as backend and can be persistent or not (in-memory).

```
noms serve ldb:<database>
```

```
noms serve mem
```

Querying datasets

List of available datasets in a database

```
noms ds ldb:<database>
```

```
noms ds http://<server:port>
```

Syncing datasets within local and remote databases

Similar to git push / pull

```
noms sync <origin> <destination>
```

Pulling from remote to local

```
noms sync http://www.acme.com:8000::ds ldb:db::ds
```

Pushing from local to remote

```
noms sync ldb:db::ds http://www.acme.com:8000::ds
```


Historical commit log of a dataset

```
noms log http://<server:port>::<dataset>
```

```
commit nb28vpmt3u0of37u1hs2lis5nqo7942j  
Parent: 4niqn6lu9kuq6i35ovq26vqfaov60vs3  
(root) {  
+   {  
+     name: "Gustavo",  
+     phone: "9876-5432",  
+   }  
}
```

```
commit 4niqn6lu9kuq6i35ovq26vqfaov60vs3  
Parent: None
```

Displaying meta data and data values

```
noms show http://<server:port>::[<dataset>][#<commit>]
```

```
struct Commit {  
  meta: struct {},  
  parents: Set<Ref<Cycle<0>>>,  
  value: List<struct {  
    given: String,  
    male: Bool,  
  }>,  
}({  
  meta: {},  
  parents: {},  
  value: [ // 4 items  
    {  
      given: "Rickon",  
      male: true,  
    },  
    {  
      given: "Bran",  
      male: true,  
    },  
    {  
      given: "Arya",  
      male: false,  
    },  
  ],  
})
```

Diffing two commits

```
noms diff http://<server:port>::#<commit-1> http://<server:port>::#<commit-2>
```

```
.parents {  
- nb28vpmt3u0of37u1hs2lis5nqo7942j  
+ 4niqn6lu9kuq6i35ovq26vqfaov60vs3  
}  
.value {  
- {  
-   name: "Luiz",  
-   phone: "1234-5678",  
- }  
}
```

JavaScript SDK

Allows a developer to work with databases, datasets, commits, data, etc.

npm install @attic/noms

<https://github.com/attic-labs/noms/blob/master/doc/js-tour.md>

DEMO



References

- [Slides / Sample code](#)
- [Noms @ GitHub](#)
- [Introduction to Noms](#)
- [JavaScript API](#)
- [NPM API Package](#)
- [Slack Team](#)
- [Attic.io](#)

Thanks!!

Gustavo Coelho

Software Engineer @ Hewlett Packard Enterprise

Email: gcoelho@gcoelho.com

Twitter: [@gafcoelho](https://twitter.com/gafcoelho)

GitHub: github.com/gcoelho

Luiz Svoboda

Software Engineer @ Hewlett Packard Enterprise

Email: luizek@gmail.com

Twitter: [@luizsvoboda](https://twitter.com/luizsvoboda)

GitHub: github.com/lhsvobodaj